

Introduction to R for Statistics & Data Science*

Peter Martey Addo

*Data Scientist de DIVISION INFORMATION & INNOVATION
SNCF - DIRECTION DU MATERIEL*

14 Février 2017

Contents

Introduction to R	1
Overview	1
Commenting and R Environments	2
Data Imports, Preparation & Manipulation	6
Getting Started and Resource	6
How To Create A Simple Data Frame in R	6
Faster Data Manipulation in R using dplyr	8
Introduction	8
filter: Keep rows matching criteria	10
select: Pick columns by name	11
Chaining Method (%>%: “then” operator)	12
arrange: Reorder rows	13
mutate: Add new variables	14
summarise: Reduce variables to values	14
Window Functions	17
Some Useful Functions	18
Descriptive Statistics, EDA & Statistical testing	20
Descriptive Statistics	20
Time Series Analysis	30
Hypothesis Testing	34
Time Series of difference of log CAC 40	35
Autocorrelation Function Plots	35
Modeling	36
Introduction to Data Science Algorithms in R	38
Ensemble Feature Selection on Titanic Dataset	38
References (Some useful R links)	40

Introduction to R

Overview

R is a powerful **open-source statistical programming language** software program for advanced data techniques. It allows you to work on:

- statistical analysis

*Diffusion Limitée



Figure 1:

- machine learning & deep learning
- data visualization.

Commenting and R Environments

Adding Comments to R Code

Comments are meant to help describe the code, and should be used when the purpose of a line of code is unclear. In R, lines beginning with a pound character (#) are each a **comment**. Comments should be clear, concise, and helpful.

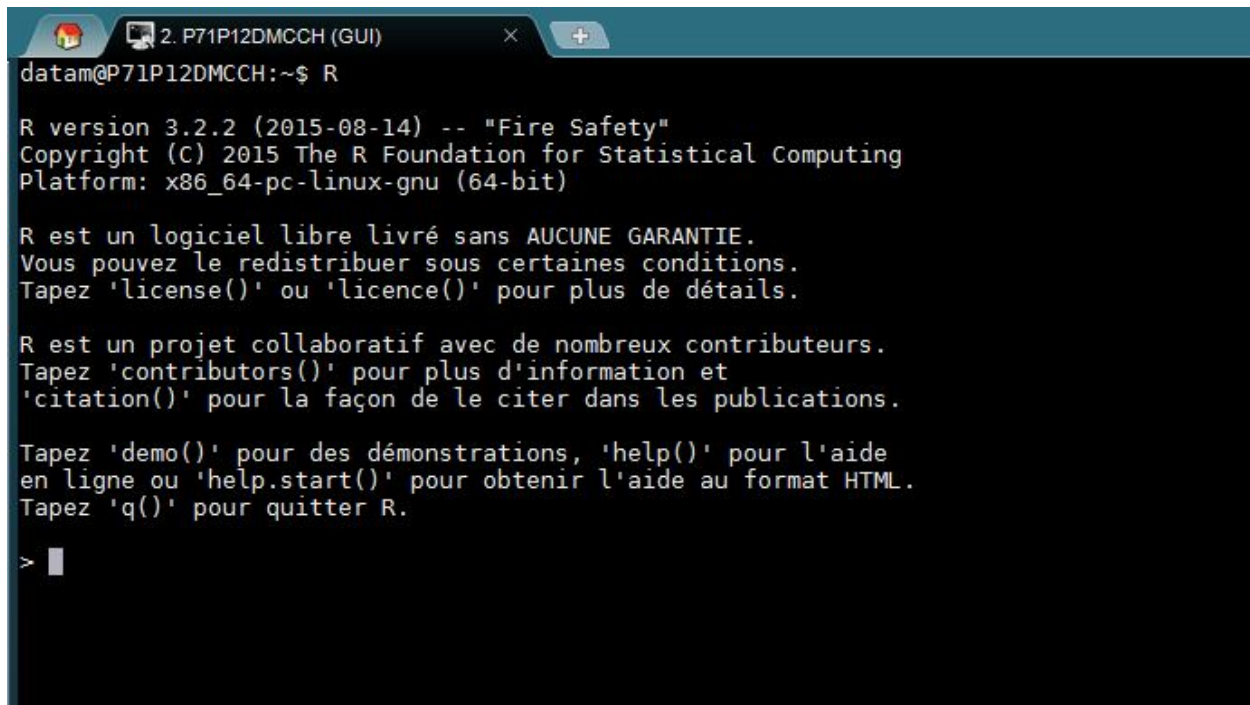
Here's an example of some commented R code, which will be explained below.

```
# Describe purpose of computation
amount.per.day <- 10 # amount in euros
result <- 10 * 7 # budget for the week
```

Interactive R Sessions

An interactive R session can be opened within terminal. To open, type the letter **R** into your terminal and hit **enter**, an interactive R session will appear :

Limitation: this environment doesn't provide much of an interface for writing R scripts.



```
datam@P71P12DMCCH:~$ R
R version 3.2.2 (2015-08-14) -- "Fire Safety"
Copyright (C) 2015 The R Foundation for Statistical Computing
Platform: x86_64-pc-linux-gnu (64-bit)

R est un logiciel libre livré sans AUCUNE GARANTIE.
Vous pouvez le redistribuer sous certaines conditions.
Tapez 'license()' ou 'licence()' pour plus de détails.

R est un projet collaboratif avec de nombreux contributeurs.
Tapez 'contributors()' pour plus d'information et
'citation()' pour la façon de le citer dans les publications.

Tapez 'demo()' pour des démonstrations, 'help()' pour l'aide
en ligne ou 'help.start()' pour obtenir l'aide au format HTML.
Tapez 'q()' pour quitter R.

> █
```

Figure 2: screenshot of interactive r session

RStudio

RStudio is an open-source **integrated development environment**(IDE) that provides an informative user-interface for interacting with the R software program (You can download the free version).

The RStudio Interface:

- **Script:** Text editor for writing your R code.
- **Console:** A console for entering R commands. This is very similar to your command-line. You can **use the up arrow** to easily access previously executed lines of code.
- **Environment:** displays information that you have stored inside of variables.
- **Files, Plots, Packages, Help etc.:** there are multiple tabs for accessing various information.

Remark: See [rstudio-IDE-cheatsheet](#) for detailed information on the IDE.

Creating Variables

In R, variable names can contain any combination of alphanumeric characters, as well as periods (.) or underscores (_) but *cannot* begin with a number, period, or underscore – they are also case sensitive.

When you are declaring a variable in R, you use the assignment operator <- to store information in a variable. For example:

```
# Stores the number 80 into a variable called monthly.transport.fee
monthly.transport.fee <- 80
```

Type the variable name into the R console and hit enter to see the information stored in a variable:

```
monthly.transport.fee
## [1] 80
```

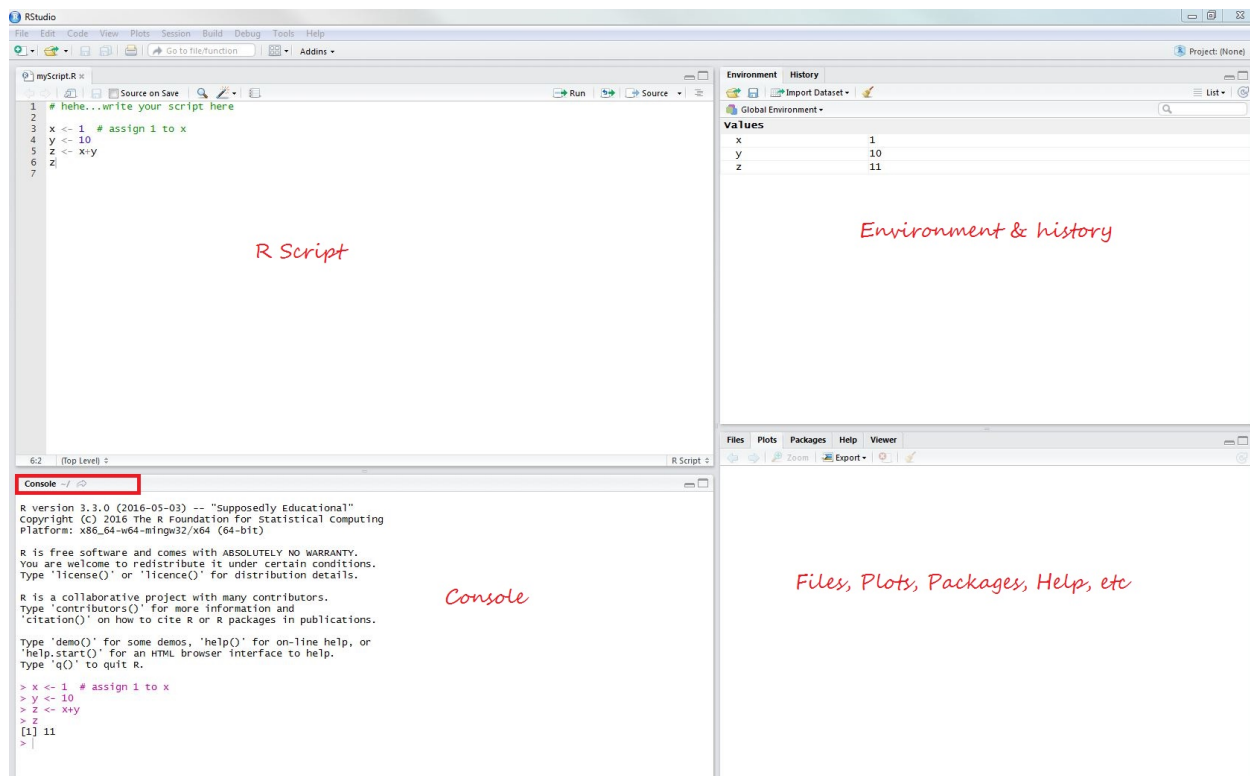


Figure 3: screenshot of labeled RStudio interface

You can also use print function

```
print(monthly.transport.fee)
## [1] 80
```

It's also quite easy to use **basic mathematical operators** (+, -, /, *) in your creation of variables. For example, you could create a variable that is the sum of two numbers as follows:

```
x <- 10 + 8
```

Other arithmetic operators include:

- Exponentiation: ^ : for example 3^2 equals 9.
- Modulo: %% : for example 5 modulo 3 or 5 %% 3 equals 2.

Workspace

- Use `ls()` to show contents in the workspace.
- Use `rm()` to remove contents in workspace.

```
# Clear the entire workspace
rm(list = ls())
```

```
# Create the variables milk, sugar and bread
milk <- 3
sugar <- 7
bread <- 1
```

```
# Inspect the contents of the workspace again
ls()
```

```
## [1] "bread" "milk" "sugar"
```

Basic Data Types

Basic data types in R:

- **Numeric:** Numeric data consists of the set of real numbers.

```
x <- 10.23 + 8
```

- **Character:** Character data stores strings of characters in a variable. Encapsulate the strings in either single (') or double quotes ("):

```
# Create character variable `name.manager` with the value "Guillaume THERY"
name.manager <- "Guillaume THERY"
```

- **Boolean:** Boolean (logical) data types can only take on two values: TRUE and FALSE.

```
# Test if 3.15 is greater than pi, and store the results in a variable `x`
x <- 3.15 > pi # returns the boolean value TRUE since `pi=3.141593`
x
##[1] TRUE
```

- **Complex:** Complex numbers are created using the i syntax:

```
complex.variable <- 4+5i # create a complex number with real=4 & imaginary=5
```

- **Integer:** Create an integer by placing a capital L after an integer value in variable assignment

```
x.integer <- 4L
```

Package Management and Installation

Packages are used to extend the functionality of R

- Installing packages from CRAN:

- One-by-one

```
install.packages("ggplot2")
```

- Group

```
list.of.packages <- c("ggplot2","reshape2","plyr","data.table","dplyr","lubridate")
new.packages <- list.of.packages[!(list.of.packages %in% installed.packages()[,"Package"])]
if(length(new.packages)) install.packages(new.packages)
```

- Installing from Github:

```
require(devtools)
install_github('ramnathv/rCharts')
```

- To load a package (need to load before using)

```
library("ggplot2")
```

- `all.packages <- available.packages()` to get the list of all the packages in CRAN

Data Imports, Preparation & Manipulation

Getting Started and Resource

```
?mean # help on a function
help.search('kurtosis') # search help files for a word
help(package = 'ggplot2') # find help in package
search() # to list all of the functions now available to you
getwd() # to get the working directory
setwd() # to change the working directory
```

In getting started with R, RHeatSheet's will be provided as supporting documents to avoid repetitions in this training session. Participants can download such reference cards via RStudio website.

Remark: We will use the **Base R CheatSheet** to learn basic R syntax: Vectors, Matrices, Lists, Data Frames, Functions and etc.

How To Create A Simple Data Frame in R

```
Died.At <- c(22,40,72,41)
Writer.At <- c(16, 18, 36, 36)
First.Name <- I(c("Kevin", "Helen", "Walt", "Janet"))
Second.Name <- I(c("Doe", "Danse", "Whitman", "Austen"))
Sex <- c("MALE", "FEMALE", "MALE", "FEMALE")
Date.Of.Death <- as.Date(c("2015-05-10", "1949-10-07", "1992-03-26", "1887-07-18"))
writers_df <- data.frame(Died.At, Writer.At, First.Name, Second.Name, Sex, Date.Of.Death)
head(writers_df) # to get the head in the data frame
```

```
##   Died.At Writer.At First.Name Second.Name    Sex Date.Of.Death
## 1      22        16      Kevin         Doe  MALE   2015-05-10
## 2      40        18      Helen         Danse FEMALE 1949-10-07
## 3      72        36      Walt         Whitman MALE   1992-03-26
## 4      41        36      Janet         Austen FEMALE 1887-07-18
```

```
tail(writers_df) # to get the tail in the data frame
```

```
##   Died.At Writer.At First.Name Second.Name    Sex Date.Of.Death
## 1      22        16      Kevin         Doe  MALE   2015-05-10
## 2      40        18      Helen         Danse FEMALE 1949-10-07
## 3      72        36      Walt         Whitman MALE   1992-03-26
## 4      41        36      Janet         Austen FEMALE 1887-07-18
```

Convert column Date.Of.Death as date format:

```
suppressMessages(library(lubridate))
```

```
## Warning: package 'lubridate' was built under R version 3.3.2
```

```
writers_df$Date.Of.Death <- ymd(writers_df$Date.Of.Death)
str(writers_df)
```

```
## 'data.frame':   4 obs. of  6 variables:
##  $ Died.At      : num  22 40 72 41
##  $ Writer.At    : num  16 18 36 36
##  $ First.Name   :Class 'AsIs' chr [1:4] "Kevin" "Helen" "Walt" "Janet"
```



```
## $ Second.Name :Class 'AsIs' chr [1:4] "Doe" "Danse" "Whitman" "Austen"
## $ Sex          : Factor w/ 2 levels "FEMALE","MALE": 2 1 2 1
## $ Date.Of.Death: Date, format: "2015-05-10" "1949-10-07" ...
```

What you need to know about Data Frames

We make use of the Airline Dataset. The data comes originally from RITA where it is described in detail.

```
data <- read.csv("allyears2k.csv") # to read a CSV file into a data frame
names(data) # get names of variables / fields in the data frame
```

```
## [1] "Year"           "Month"           "DayofMonth"
## [4] "DayOfWeek"      "DepTime"          "CRSDepTime"
## [7] "ArrTime"        "CRSArrTime"       "UniqueCarrier"
## [10] "FlightNum"      "TailNum"          "ActualElapsedTime"
## [13] "CRSElapsedTime" "AirTime"          "ArrDelay"
## [16] "DepDelay"       "Origin"           "Dest"
## [19] "Distance"       "TaxiIn"           "TaxiOut"
## [22] "Cancelled"      "CancellationCode" "Diverted"
## [25] "CarrierDelay"   "WeatherDelay"     "NASDelay"
## [28] "SecurityDelay"  "LateAircraftDelay" "IsArrDelayed"
## [31] "IsDepDelayed"
```

```
dim(data) # dimension of data frame
```

```
## [1] 43978 31
```

```
nrow(data) # to get the number of rows in the data frame
```

```
## [1] 43978
```

```
ncol(data) # to get the number of columns in the data frame
```

```
## [1] 31
```

```
head(data,2) # to get the head of 2 rows in the data frame
```

```
##   Year Month DayofMonth DayOfWeek DepTime CRSDepTime ArrTime CRSArrTime
## 1 1987   10         14         3    741      730      912      849
## 2 1987   10         15         4    729      730      903      849
##   UniqueCarrier FlightNum TailNum ActualElapsedTime CRSElapsedTime AirTime
## 1             PS     1451    <NA>             91             79      NA
## 2             PS     1451    <NA>             94             79      NA
##   ArrDelay DepDelay Origin Dest Distance TaxiIn TaxiOut Cancelled
## 1       23       11   SAN  SFO     447     NA     NA        0
## 2       14       -1   SAN  SFO     447     NA     NA        0
##   CancellationCode Diverted CarrierDelay WeatherDelay NASDelay
## 1             <NA>        0         NA         NA         NA
## 2             <NA>        0         NA         NA         NA
##   SecurityDelay LateAircraftDelay IsArrDelayed IsDepDelayed
## 1             NA                 NA         YES         YES
## 2             NA                 NA         YES         NO
```

```
tail(data,2) # to get the tail of 2 rows in the data frame
```

```
##   Year Month DayofMonth DayOfWeek DepTime CRSDepTime ArrTime
## 43977 2008     1         3         4   1652      1625      1750
## 43978 2008     1         3         4   1057      1030      1154
```

```
##      CRSArrTime UniqueCarrier FlightNum TailNum ActualElapsedTime
## 43977      1730           WN      1022  N655WN              58
## 43978      1140           WN      1041  N474WN              57
##      CRSElapsedTime AirTime ArrDelay DepDelay Origin Dest Distance TaxiIn
## 43977           65      44      20      27   SLC  BOI      291      4
## 43978           70      46      14      27   SLC  BOI      291      2
##      TaxiOut Cancelled CancellationCode Diverted CarrierDelay
## 43977          10          0              0              0
## 43978           9          0              0              NA
##      WeatherDelay NASDelay SecurityDelay LateAircraftDelay IsArrDelayed
## 43977           0          4              0              16      YES
## 43978          NA          NA              NA              NA      YES
##      IsDepDelayed
## 43977          YES
## 43978          YES
```

```
str(data) # to get structure of the data frame
attributes(data) # to see the data frame attributes
summary(data) # to see data summary
data$Year # to see the column Year...same as data[[1]]
data[1] # to create a new data frame with just column 1
data[1:4] # to create a data frame with first four columns
data[, 1] # to get column 1 listed
data$Year[20] # to get the 20th element in the Year column
tail(data,2) # to get the tail of 2 rows in the data frame
write.csv(data,"mydataname.csv", row.names = FALSE) # save a dataset in csv
```

- thus data[,1][47] is then the same as data\$Year[47]
- length(data\$IsArrDelayed[is.na(data\$IsArrDelayed)]) to find the number of NAs in a given column (IsArrDelayed) in a data frame

```
# to get the mean of the valid numbers in the group
mean(data$Cancelled[!is.na(data$Cancelled)])
```

```
## [1] 0.02469417
```

```
# to get the mean value of Departure Delay where Arrival Delay is > 10 and Day Of Week is Monday
mean(subset(data, !is.na(data$CarrierDelay) & data$ArrDelay > 10 & data$DayOfWeek == 1)$DepDelay)
```

```
## [1] 45.98084
```

Faster Data Manipulation in R using dplyr

Introduction

- Five basic verbs: filter, select, arrange, mutate, summarise (plus group_by)
- dplyr will mask a few base functions
- The data on Régularité mensuelle TGV depuis septembre 2011. The data can be download from the SNCF Open data platform. Click on download for data.

```
# load packages
suppressMessages(library(dplyr))
```

```
## Warning: package 'dplyr' was built under R version 3.3.2
```



```

suppressMessages(library(lubridate))

# import data from current working directory (eg data in the folder "opendatasncf")

regularite.mensuelle.tgv <- read.csv("opendatasncf/regularite-mensuelle-tgv.csv",
                                     encoding="UTF-8", sep=";")

# view data

head(regularite.mensuelle.tgv)

```

```

##      Date      Axe      Départ      Arrivée
## 1 2014-08 Atlantique  ANGERS SAINT LAUD  PARIS MONTPARNASSE
## 2 2014-08 Atlantique  ANGOULEME  PARIS MONTPARNASSE
## 3 2014-08 Sud-Est  ANNECY  PARIS LYON
## 4 2014-08 Sud-Est  AVIGNON TGV  PARIS LYON
## 5 2014-08 Sud-Est BESANCON FRANCHE COMTE TGV  PARIS LYON
## 6 2014-08 Nord  LILLE MARSEILLE ST CHARLES
##  Nombre.de.trains.programmés  Nombre.de.trains.ayant.circulé
## 1 440 440
## 2 327 327
## 3 124 124
## 4 524 524
## 5 214 214
## 6 124 124
##  Nombre.de.trains.annulés  Nombre.de.trains.en.retard.à.l.arrivée
## 1 0 21
## 2 0 31
## 3 0 5
## 4 0 89
## 5 0 9
## 6 0 13
## Régularité
## 1 95.2
## 2 90.5
## 3 96.0
## 4 83.0
## 5 95.8
## 6 89.5

```

```

##
## 1
## 2
## 3
## 4 En août, la régularité des TGV entre Paris et la Méditerranée a été fortement impactée par 8 acciden
## 5
## 6

```

- `tbl_df` creates a “local data frame”
- Local data frame is simply a wrapper for a data frame that prints nicely

```

# convert to local data frame

regularite <- tbl_df(regularite.mensuelle.tgv)

# rename columns by removing french accents

```

```
colnames(regularite)<-c("Date","Axe","Depart","Arrivee","Nombre.de.trains.programmes",
  "Nombre.de.trains.ayant.circule","Nombre.de.trains.annules",
  "Nombre.de.trains.en.retard.a.l.arrivee","Regularite","Commentaires")

# printing only shows 10 rows and as many columns as can fit on your screen
regularite

## # A tibble: 5,900 × 10
##   Date      Axe      Depart      Arrivee
##   <fctr>    <fctr>    <fctr>    <fctr>
## 1 2014-08 Atlantique ANGERS SAINT LAUD PARIS MONTPARNASSE
## 2 2014-08 Atlantique ANGOULEME PARIS MONTPARNASSE
## 3 2014-08 Sud-Est ANNECY PARIS LYON
## 4 2014-08 Sud-Est AVIGNON TGV PARIS LYON
## 5 2014-08 Sud-Est BESANCON FRANCHE COMTE TGV PARIS LYON
## 6 2014-08 Nord LILLE MARSEILLE ST CHARLES
## 7 2014-08 Sud-Est LYON PART DIEU MONTPELLIER
## 8 2014-08 Sud-Est MACON LOCHE PARIS LYON
## 9 2014-08 Sud-Est MULHOUSE VILLE PARIS LYON
## 10 2014-08 Est PARIS EST STRASBOURG
## # ... with 5,890 more rows, and 6 more variables:
## #   Nombre.de.trains.programmes <int>,
## #   Nombre.de.trains.ayant.circule <int>, Nombre.de.trains.annules <int>,
## #   Nombre.de.trains.en.retard.a.l.arrivee <int>, Regularite <dbl>,
## #   Commentaires <fctr>

# you can change Date in date format
regularite$Date1 <- paste0(regularite$Date,sep='-',01) # add a day part of date
regularite$Date2 <- ymd(regularite$Date1) # date column
regularite$Year <- year(regularite$Date2)
regularite$Month <- month(regularite$Date2)
regularite$Quarters <- quarters(regularite$Date2)

# you can specify that you want to see more rows
print(regularite, n=20)

# convert to a normal data frame to see all of the columns
data.frame(head(regularite))
```

filter: Keep rows matching criteria

- Base R approach to filtering forces you to repeat the data frame's name
- dplyr approach is simpler to write and read
- Command structure (for all dplyr verbs):
 - first argument is a data frame
 - return value is a data frame
 - nothing is modified in place

```
# base R approach to view all regularite on Quarter 1 and January
regularite[regularite$Quarters=='Q1' & regularite$Month==1, ]
#another method
subset(regularite,regularite$Quarters=='Q1' & regularite$Month==1)
subset(regularite,Quarters=='Q1' & Month==1)
```

```

# dplyr approach
# note: you can use comma or ampersand to represent AND condition
filter(regularite, Quarters=='Q1', Month==1)

## # A tibble: 500 × 15
##   Date      Axe      Depart      Arrivee
##   <fctr> <fctr>      <fctr>      <fctr>
## 1 2015-01 Nord      DOUAI      PARIS NORD
## 2 2015-01 Sud-Est  MACON LOCHE  PARIS LYON
## 3 2015-01 Sud-Est MARSEILLE ST CHARLES  LYON PART DIEU
## 4 2015-01 Est      NANCY      PARIS EST
## 5 2015-01 Sud-Est  NICE VILLE  PARIS LYON
## 6 2015-01 Sud-Est  PARIS LYON BESANCON FRANCHE COMTE TGV
## 7 2012-01 Sud-Est  AIX EN PROVENCE TGV  PARIS LYON
## 8 2012-01 Sud-Est  DIJON VILLE  PARIS LYON
## 9 2012-01 Nord      PARIS NORD  DOUAI
## 10 2012-01 Sud-Est  PARIS LYON  GRENOBLE
## # ... with 490 more rows, and 11 more variables:
## #   Nombre.de.trains.programmes <int>,
## #   Nombre.de.trains.ayant.circule <int>, Nombre.de.trains.annules <int>,
## #   Nombre.de.trains.en.retard.a.l.arrivee <int>, Regularite <dbl>,
## #   Commentaires <fctr>, Date1 <chr>, Date2 <date>, Year <dbl>,
## #   Month <dbl>, Quarters <chr>

# use pipe for OR condition
filter(regularite, Depart=="LYON PART DIEU" | Depart=="PARIS NORD")

## # A tibble: 531 × 15
##   Date      Axe      Depart      Arrivee
##   <fctr> <fctr>      <fctr>      <fctr>
## 1 2014-08 Sud-Est LYON PART DIEU  MONTPELLIER
## 2 2014-10 Sud-Est LYON PART DIEU MARSEILLE ST CHARLES
## 3 2014-10 Nord      PARIS NORD  DOUAI
## 4 2014-12 Nord      PARIS NORD  DOUAI
## 5 2015-06 Nord      LYON PART DIEU  LILLE
## 6 2015-07 Nord      PARIS NORD  DOUAI
## 7 2015-07 Nord      PARIS NORD  DUNKERQUE
## 8 2015-08 Sud-Est LYON PART DIEU MARSEILLE ST CHARLES
## 9 2015-08 Nord      PARIS NORD  DOUAI
## 10 2015-09 Nord      PARIS NORD  DOUAI
## # ... with 521 more rows, and 11 more variables:
## #   Nombre.de.trains.programmes <int>,
## #   Nombre.de.trains.ayant.circule <int>, Nombre.de.trains.annules <int>,
## #   Nombre.de.trains.en.retard.a.l.arrivee <int>, Regularite <dbl>,
## #   Commentaires <fctr>, Date1 <chr>, Date2 <date>, Year <dbl>,
## #   Month <dbl>, Quarters <chr>

# you can also use %in% operator
filter(regularite, Depart %in% c("LYON PART DIEU", "PARIS NORD"))

```

select: Pick columns by name

- Base R approach is awkward to type and to read
- dplyr approach uses similar syntax to `filter`

- Like a SELECT in SQL

```
# base R approach to select Depart, Arrivee, and Nombre.de.trains.programmes columns
regularite[, c("Depart", "Arrivee", "Nombre.de.trains.programmes")]
```

```
# dplyr approach
```

```
select(regularite, Depart, Arrivee, Nombre.de.trains.programmes)
```

```
## # A tibble: 5,900 × 3
```

```
##           Depart           Arrivee
##           <fctr>          <fctr>
## 1     ANGERS SAINT LAUD  PARIS MONTPARNASSE
## 2           ANGOULEME  PARIS MONTPARNASSE
## 3           ANNECY      PARIS LYON
## 4     AVIGNON TGV      PARIS LYON
## 5  BESANCON FRANCHE COMTE TGV  PARIS LYON
## 6           LILLE MARSEILLE ST CHARLES
## 7     LYON PART DIEU  MONTPELLIER
## 8           MACON LOCHE  PARIS LYON
## 9     MULHOUSE VILLE  PARIS LYON
## 10          PARIS EST    STRASBOURG
## # ... with 5,890 more rows, and 1 more variables:
## #   Nombre.de.trains.programmes <int>
```

```
# use colon to select multiple contiguous columns, and use `contains` to match columns by name
# note: `starts_with`, `ends_with`, and `matches` (for regular expressions)
```

```
# can also be used to match columns by name
```

```
select(regularite, Year:Month, contains("annules"), contains("retard"))
```

```
## # A tibble: 5,900 × 4
```

```
##   Year Month Nombre.de.trains.annules
##   <dbl> <dbl>          <int>
## 1  2014     8              0
## 2  2014     8              0
## 3  2014     8              0
## 4  2014     8              0
## 5  2014     8              0
## 6  2014     8              0
## 7  2014     8              0
## 8  2014     8              0
## 9  2014     8              0
## 10 2014     8              0
```

```
## # ... with 5,890 more rows, and 1 more variables:
```

```
## #   Nombre.de.trains.en.retard.a.l.arrivee <int>
```

Chaining Method (%>%: “then” operator)

- Usual way to perform multiple operations in one line is by nesting
- Chaining increases readability significantly when there are many commands

```
# nesting method to select Depart and Nombre.de.trains.en.retard.a.l.arrivee columns
# and filter for delays over 60 minutes
```

```
filter(select(regularite, Depart, Nombre.de.trains.en.retard.a.l.arrivee),
        Nombre.de.trains.en.retard.a.l.arrivee > 60)
```

```
# chaining method
regularite %>%
  select(Depart, Nombre.de.trains.en.retard.a.l.arrivee) %>%
  filter(Nombre.de.trains.en.retard.a.l.arrivee > 60)

## # A tibble: 662 × 2
##       Depart Nombre.de.trains.en.retard.a.l.arrivee
##       <fctr>                                     <int>
## 1     AVIGNON TGV                                89
## 2     LYON PART DIEU                             75
## 3  MARSEILLE ST CHARLES                          105
## 4     LYON PART DIEU                             147
## 5    AIX EN PROVENCE TGV                           69
## 6     AVIGNON TGV                                77
## 7  MARSEILLE ST CHARLES                          155
## 8     MONTPELLIER                                79
## 9    ST PIERRE DES CORPS                          73
## 10  MARSEILLE ST CHARLES                          79
## # ... with 652 more rows
```

- Operator is automatically imported from the magrittr package
- Can be used to replace nesting in R commands outside of dplyr

```
# create two vectors and calculate Euclidian distance between them
a <- 3:7; b <- 1:5
sqrt(sum((a-b)^2))
```

```
# chaining method
(a-b)^2 %>% sum() %>% sqrt()
```

```
## [1] 4.472136
```

arrange: Reorder rows

```
# base R approach to select Depart and Nombre.de.trains.en.retard.a.l.arrivee columns and sort by
# Nombre.de.trains.en.retard.a.l.arrivee
regularite[order(regularite$Nombre.de.trains.en.retard.a.l.arrivee),
  c("Depart", "Nombre.de.trains.en.retard.a.l.arrivee")]
```

```
# dplyr approach
regularite %>%
  select(Depart, Nombre.de.trains.en.retard.a.l.arrivee) %>%
  arrange(Nombre.de.trains.en.retard.a.l.arrivee)
```

```
## # A tibble: 5,900 × 2
##       Depart Nombre.de.trains.en.retard.a.l.arrivee
##       <fctr>                                     <int>
## 1     PARIS LYON                                0
## 2  SAINT ETIENNE CHATEAUCREUX                    0
## 3     PARIS MONTPARNASSE                          0
## 4  SAINT ETIENNE CHATEAUCREUX                    0
## 5     PARIS LYON                                0
## 6  SAINT ETIENNE CHATEAUCREUX                    0
## 7  SAINT ETIENNE CHATEAUCREUX                    0
## 8     PARIS LYON                                0
```

```
## 9          PARIS LYON          0
## 10         PARIS LYON          0
## # ... with 5,890 more rows
```

```
# use `desc` for descending
regularite %>%
  select(Depart, Nombre.de.trains.en.retard.a.l.arrivee) %>%
  arrange(desc(Nombre.de.trains.en.retard.a.l.arrivee))
```

mutate: Add new variables

- Create new variables that are functions of existing variables

```
# base R approach to create a new variable Non.Regularite (in %)
regularite$Non.Regularite <- 100-regularite$Regularite
regularite[, c("Regularite", "Arrivee", "Non.Regularite")]
```

```
# dplyr approach (prints the new variable but does not store it)
regularite %>%
  select(Regularite, Arrivee) %>%
  mutate(Non.Regularite = 100-Regularite)
```

```
## # A tibble: 5,900 × 3
##   Regularite      Arrivee Non.Regularite
##   <dbl>         <fctr>         <dbl>
## 1    95.2    PARIS MONT-PARNASSE      4.8
## 2    90.5    PARIS MONT-PARNASSE      9.5
## 3    96.0      PARIS LYON          4.0
## 4    83.0      PARIS LYON         17.0
## 5    95.8      PARIS LYON          4.2
## 6    89.5 MARSEILLE ST CHARLES     10.5
## 7    81.1      MONTPELLIER     18.9
## 8    86.7      PARIS LYON         13.3
## 9    93.4      PARIS LYON          6.6
## 10   93.6      STRASBOURG         6.4
## # ... with 5,890 more rows
```

```
# store the new variable
regularite <- regularite %>% mutate(Non.Regularite = 100-Regularite)
```

summarise: Reduce variables to values

- Useful with data that has been grouped by one or more variables
- `group_by` creates the groups that will be operated on
- `summarise` uses the provided aggregation function to summarise each group

```
# base R approaches to calculate the average arrival delay to each destination
head(with(regularite, tapply(Nombre.de.trains.en.retard.a.l.arrivee, Arrivee, mean, na.rm=TRUE)))
head(aggregate(Nombre.de.trains.en.retard.a.l.arrivee ~ Arrivee, regularite, mean))
```

```
# dplyr approach: create a table grouped by Arrivee, and then summarise each group
# by taking the mean of Nombre.de.trains.en.retard.a.l.arrivee
regularite %>%
  group_by(Arrivee) %>%
  summarise(avg_delay = mean(Nombre.de.trains.en.retard.a.l.arrivee, na.rm=TRUE))
```

```
## # A tibble: 48 × 2
##               Arrivee avg_delay
##               <fctr>    <dbl>
## 1      AIX EN PROVENCE TGV  43.54237
## 2      ANGERS SAINT LAUD  31.03390
## 3      ANGOULEME        23.44068
## 4      ANNECY           11.88136
## 5      ARRAS            31.11864
## 6      AVIGNON TGV       40.94915
## 7      BELLEGARDE (AIN)  30.47458
## 8      BESANCON FRANCHE COMTE TGV 15.79661
## 9      BORDEAUX ST JEAN  50.30508
## 10     BREST           11.52542
## # ... with 38 more rows
```

- summarise_each allows you to apply the same summary function to multiple columns at once
- Note: mutate_each is also available

```
# for each Depart, calculate the mean of Nombre.de.trains.annules or
# Nombre.de.trains.en.retard.a.l.arrivee
regularite %>%
  group_by(Depart) %>%
  summarise_each(funs(mean), Nombre.de.trains.annules,
                 Nombre.de.trains.en.retard.a.l.arrivee)
```

```
## # A tibble: 48 × 3
##               Depart Nombre.de.trains.annules
##               <fctr>                <dbl>
## 1      AIX EN PROVENCE TGV              0.2711864
## 2      ANGERS SAINT LAUD              2.0508475
## 3      ANGOULEME                    2.1525424
## 4      ANNECY                      1.0847458
## 5      ARRAS                      1.9661017
## 6      AVIGNON TGV                  1.9830508
## 7      BELLEGARDE (AIN)             0.8135593
## 8      BESANCON FRANCHE COMTE TGV    0.7457627
## 9      BORDEAUX ST JEAN             3.3050847
## 10     BREST                      0.4915254
## # ... with 38 more rows, and 1 more variables:
## #   Nombre.de.trains.en.retard.a.l.arrivee <dbl>
```

```
# for each Depart, calculate the minimum and maximum arrival delays
regularite %>%
  group_by(Depart) %>%
  summarise_each(funs(min(., na.rm=TRUE), max(., na.rm=TRUE)),
                 matches("Nombre.de.trains.en.retard.a.l.arrivee"))
```

```
## # A tibble: 48 × 3
##               Depart    min    max
##               <fctr> <int> <int>
## 1      AIX EN PROVENCE TGV      0     95
## 2      ANGERS SAINT LAUD     16    115
## 3      ANGOULEME           14    102
## 4      ANNECY              5     24
## 5      ARRAS              19    109
```



```
## 6          AVIGNON TGV      24    140
## 7          BELLEGARDE (AIN)    6     84
## 8  BESANCON FRANCHE COMTE TGV    2     32
## 9          BORDEAUX ST JEAN   34    140
## 10         BREST            2     29
## # ... with 38 more rows
```

- Helper function `n()` counts the number of rows in a group
- Helper function `n_distinct(vector)` counts the number of unique items in that vector

```
# for each month of the year, count the total number of rows and sort in descending order
regularite %>%
```

```
  group_by(Year, Month) %>%
  summarise(row_count = n()) %>%
  arrange(desc(row_count))
```

```
## Source: local data frame [59 x 3]
```

```
## Groups: Year [6]
```

```
##
##   Year Month row_count
##   <dbl> <dbl>   <int>
## 1  2011     9       100
## 2  2011    10       100
## 3  2011    11       100
## 4  2011    12       100
## 5  2012     1       100
## 6  2012     2       100
## 7  2012     3       100
## 8  2012     4       100
## 9  2012     5       100
## 10 2012     6       100
## # ... with 49 more rows
```

```
# rewrite more simply with the `tally` function
```

```
regularite %>%
  group_by(Year, Month) %>%
  tally(sort = TRUE)
```

```
## Source: local data frame [59 x 3]
```

```
## Groups: Year [6]
```

```
##
##   Year Month    n
##   <dbl> <dbl> <int>
## 1  2011     9   100
## 2  2011    10   100
## 3  2011    11   100
## 4  2011    12   100
## 5  2012     1   100
## 6  2012     2   100
## 7  2012     3   100
## 8  2012     4   100
## 9  2012     5   100
## 10 2012     6   100
## # ... with 49 more rows
```

```
# for each Axe, count the total number of distinct destination
regularite %>%
  group_by(Axe) %>%
  summarise(Arrivee_count = n_distinct(Arrivee))
```

```
## # A tibble: 4 × 2
##       Axe Arrivee_count
##   <fctr>      <int>
## 1 Atlantique      18
## 2 Est             8
## 3 Nord            7
## 4 Sud-Est        21
```

Window Functions

- Aggregation function (like `mean`) takes `n` inputs and returns 1 value
- Window function takes `n` inputs and returns `n` values
 - Includes ranking and ordering functions (like `min_rank`), offset functions (`lead` and `lag`), and cumulative aggregates (like `cummean`).

```
# for each Departure, calculate which three months of the year they had their
# longest departure delays
# note: smallest (not largest) value is ranked as 1, so you have to
# use `desc` to rank by largest value
regularite %>%
  group_by(Depart) %>%
  select(Quarters, Month, Nombre.de.trains.en.retard.a.l.arrivee) %>%
  filter(min_rank(desc(Nombre.de.trains.en.retard.a.l.arrivee)) <= 3) %>%
  arrange(Depart, desc(Nombre.de.trains.en.retard.a.l.arrivee))
```

```
## Adding missing grouping variables: `Depart`
```

```
# rewrite more simply with the `top_n` function
regularite %>%
  group_by(Depart) %>%
  select(Quarters, Month, Nombre.de.trains.en.retard.a.l.arrivee) %>%
  top_n(3) %>%
  arrange(Depart, desc(Nombre.de.trains.en.retard.a.l.arrivee))
```

```
## Adding missing grouping variables: `Depart`
```

```
## Selecting by Nombre.de.trains.en.retard.a.l.arrivee
```

```
## Source: local data frame [151 x 4]
```

```
## Groups: Depart [48]
```

```
##
##       Depart Quarters Month
##   <fctr>    <chr> <dbl>
## 1 AIX EN PROVENCE TGV      Q3      7
## 2 AIX EN PROVENCE TGV      Q2      5
## 3 AIX EN PROVENCE TGV      Q4     11
## 4  ANGERS SAINT LAUD      Q4     11
## 5  ANGERS SAINT LAUD      Q3      7
## 6  ANGERS SAINT LAUD      Q4     10
## 7      ANGOULEME      Q1      2
## 8      ANGOULEME      Q4     11
```

```
## 9          ANGOULEME      Q3      7
## 10         ANGOULEME      Q4     10
## # ... with 141 more rows, and 1 more variables:
## #   Nombre.de.trains.en.retard.a.l.arrivee <int>

# for each Quarter, calculate the number of observations and the change from the previous Quarter
regularite %>%
  group_by(Quarters) %>%
  summarise(row_count = n()) %>%
  mutate(change = row_count - lag(row_count))

## # A tibble: 4 × 3
##   Quarters row_count change
##   <chr>     <int> <int>
## 1     Q1      1500    NA
## 2     Q2      1500     0
## 3     Q3      1400   -100
## 4     Q4      1500    100

# rewrite more simply with the `tally` function
regularite %>%
  group_by(Quarters) %>%
  tally() %>%
  mutate(change = n - lag(n))

## # A tibble: 4 × 3
##   Quarters      n change
##   <chr> <int> <int>
## 1     Q1  1500    NA
## 2     Q2  1500     0
## 3     Q3  1400   -100
## 4     Q4  1500    100
```

Some Useful Functions

```
# randomly sample a fixed number of rows, without replacement
regularite %>% sample_n(5)

## # A tibble: 5 × 16
##   Date      Axe      Depart      Arrivee
##   <fctr> <fctr>      <fctr>      <fctr>
## 1 2012-11 Sud-Est CHAMBERY CHALLES LES EAUX PARIS LYON
## 2 2014-02 Sud-Est      NICE VILLE PARIS LYON
## 3 2012-10      Est      STRASBOURG NANTES
## 4 2016-03      Est      PARIS EST NANCY
## 5 2012-07 Sud-Est      PARIS LYON AVIGNON TGV
## # ... with 12 more variables: Nombre.de.trains.programmes <int>,
## #   Nombre.de.trains.ayant.circule <int>, Nombre.de.trains.annules <int>,
## #   Nombre.de.trains.en.retard.a.l.arrivee <int>, Regularite <dbl>,
## #   Commentaires <fctr>, Date1 <chr>, Date2 <date>, Year <dbl>,
## #   Month <dbl>, Quarters <chr>, Non.Regularite <dbl>

# randomly sample a fraction of rows, with replacement
regularite %>% sample_frac(0.25, replace=TRUE)
```

```
## # A tibble: 1,475 × 16
##   Date      Axe      Depart      Arrivee
##   <fctr>    <fctr>    <fctr>    <fctr>
## 1 2015-08    Nord    MARSEILLE ST CHARLES    LILLE
## 2 2012-05    Sud-Est    PERPIGNAN    PARIS LYON
## 3 2012-05    Nord    MARSEILLE ST CHARLES    LILLE
## 4 2012-02    Sud-Est    SAINT ETIENNE CHATEAUCREUX    PARIS LYON
## 5 2014-10 Atlantique    ANGOULEME    PARIS MONTPARNASSE
## 6 2012-02    Sud-Est    NIMES    PARIS LYON
## 7 2014-01    Sud-Est    PARIS LYON AIX EN PROVENCE TGV
## 8 2013-02      Est    PARIS EST    NANCY
## 9 2013-06    Sud-Est    MARSEILLE ST CHARLES    PARIS LYON
## 10 2013-08    Sud-Est    ANNECY    PARIS LYON
## # ... with 1,465 more rows, and 12 more variables:
## #   Nombre.de.trains.programmes <int>,
## #   Nombre.de.trains.ayant.circule <int>, Nombre.de.trains.annules <int>,
## #   Nombre.de.trains.en.retard.a.l.arrivee <int>, Regularite <dbl>,
## #   Commentaires <fctr>, Date1 <chr>, Date2 <date>, Year <dbl>,
## #   Month <dbl>, Quarters <chr>, Non.Regularite <dbl>
```

```
# base R approach to view the structure of an object
str(regularite)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':   5900 obs. of  16 variables:
## $ Date      : Factor w/ 59 levels "2011-09","2011-10",...: 36 36 36 36 36 ...
## $ Axe      : Factor w/ 4 levels "Atlantique","Est",...: 1 1 4 4 4 3 4 4 ...
## $ Depart   : Factor w/ 48 levels "AIX EN PROVENCE TGV",...: 2 3 4 6 8 2 ...
## $ Arrivee  : Factor w/ 48 levels "AIX EN PROVENCE TGV",...: 33 33 32 32 ...
## $ Nombre.de.trains.programmes : int  440 327 124 524 214 124 397 188 318 440 ...
## $ Nombre.de.trains.ayant.circule : int  440 327 124 524 214 124 397 188 318 440 ...
## $ Nombre.de.trains.annules : int   0 0 0 0 0 0 0 0 0 0 ...
## $ Nombre.de.trains.en.retard.a.l.arrivee: int   21 31 5 89 9 13 75 25 21 28 ...
## $ Regularite : num  95.2 90.5 96 83 95.8 89.5 81.1 86.7 93.4 93.6 ...
## $ Commentaires : Factor w/ 437 levels "", "\"Les ralentissements pour trava...
## $ Date1      : chr  "2014-08-1" "2014-08-1" "2014-08-1" "2014-08-1" ...
## $ Date2      : Date, format: "2014-08-01" "2014-08-01" ...
## $ Year       : num  2014 2014 2014 2014 2014 ...
## $ Month      : num   8 8 8 8 8 8 8 8 8 ...
## $ Quarters   : chr   "Q3" "Q3" "Q3" "Q3" ...
## $ Non.Regularite : num  4.8 9.5 4 17 4.2 10.5 18.9 13.3 6.6 6.4 ...
```

```
# dplyr approach: better formatting, and adapts to your screen width
glimpse(regularite)
```

```
## Observations: 5,900
## Variables: 16
## $ Date      <fctr> 2014-08, 2014-08, 2014...
## $ Axe      <fctr> Atlantique, Atlantique...
## $ Depart   <fctr> ANGERS SAINT LAUD, ANG...
## $ Arrivee  <fctr> PARIS MONTPARNASSE, PA...
## $ Nombre.de.trains.programmes <int> 440, 327, 124, 524, 214...
## $ Nombre.de.trains.ayant.circule <int> 440, 327, 124, 524, 214...
## $ Nombre.de.trains.annules <int> 0, 0, 0, 0, 0, 0, 0, 0,...
## $ Nombre.de.trains.en.retard.a.l.arrivee <int> 21, 31, 5, 89, 9, 13, 7...
## $ Regularite <dbl> 95.2, 90.5, 96.0, 83.0,...
## $ Commentaires <fctr> , , , En août, la régu...
```

```
## $ Date1 <chr> "2014-08-1", "2014-08-1..."
## $ Date2 <date> 2014-08-01, 2014-08-01...
## $ Year <dbl> 2014, 2014, 2014, 2014,...
## $ Month <dbl> 8, 8, 8, 8, 8, 8, 8, 8,...
## $ Quarters <chr> "Q3", "Q3", "Q3", "Q3",...
## $ Non.Regularite <dbl> 4.8, 9.5, 4.0, 17.0, 4....
```

Descriptive Statistics, EDA & Statistical testing

Descriptive Statistics

Let us play with the Airline Dataset previously loaded in R

```
names(data) # Airline dataset
```

```
## [1] "Year"           "Month"           "DayofMonth"
## [4] "DayOfWeek"      "DepTime"         "CRSDepTime"
## [7] "ArrTime"        "CRSArrTime"      "UniqueCarrier"
## [10] "FlightNum"      "TailNum"         "ActualElapsedTime"
## [13] "CRSElapsedTime" "AirTime"         "ArrDelay"
## [16] "DepDelay"       "Origin"          "Dest"
## [19] "Distance"       "TaxiIn"          "TaxiOut"
## [22] "Cancelled"      "CancellationCode" "Diverted"
## [25] "CarrierDelay"   "WeatherDelay"    "NASDelay"
## [28] "SecurityDelay"  "LateAircraftDelay" "IsArrDelayed"
## [31] "IsDepDelayed"
```

Contingency table

```
# Create new column by grouping Departure Delay
data$groupDepDelay<- with(data,cut(DepDelay, quantile(DepDelay,na.rm=T)))
# Create contingency table
my.table0<-with(data, table(groupDepDelay,DayOfWeek))
prop.table(my.table0,1)

##           DayOfWeek
## groupDepDelay      1      2      3      4      5
## (-16,-2] 0.12530255 0.16198101 0.20350028 0.15686092 0.11320052
## (-2,1]   0.14458502 0.17561868 0.14848413 0.15548659 0.12660142
## (1,10]   0.13160541 0.16836108 0.14813987 0.18221428 0.13719138
## (10,473] 0.12296547 0.15570609 0.13801863 0.24912974 0.13322043
##           DayOfWeek
## groupDepDelay      6      7
## (-16,-2] 0.12297524 0.11617948
## (-2,1]   0.12158829 0.12763587
## (1,10]   0.10546308 0.12702491
## (10,473] 0.08918995 0.11176969

# Alternative Approach ... direct
# to build a contingency table of the counts at each combination of factor levels
with(data,table(DayOfWeek,IsArrDelayed)) # same as table(data$DayOfWeek,data$IsArrDelayed)

##           IsArrDelayed
```

```
## DayOfWeek NO YES
##      1 2559 3243
##      2 3421 3882
##      3 3285 3746
##      4 3188 4921
##      5 2250 3339
##      6 2366 2501
##      7 2468 2809

# to get probabilities of Deptime Delay on each Day of Week
my.table<-with(data, table(cut(DepDelay, quantile(DepDelay,na.rm=T)),DayOfWeek))
prop.table(my.table,1)
```

```
##      DayOfWeek
##      1      2      3      4      5
## (-16,-2] 0.12530255 0.16198101 0.20350028 0.15686092 0.11320052
## (-2,1]   0.14458502 0.17561868 0.14848413 0.15548659 0.12660142
## (1,10]   0.13160541 0.16836108 0.14813987 0.18221428 0.13719138
## (10,473] 0.12296547 0.15570609 0.13801863 0.24912974 0.13322043
##      DayOfWeek
##      6      7
## (-16,-2] 0.12297524 0.11617948
## (-2,1]   0.12158829 0.12763587
## (1,10]   0.10546308 0.12702491
## (10,473] 0.08918995 0.11176969
```

Balloon Plot

```
suppressMessages(library(gplots))

## Warning: package 'gplots' was built under R version 3.3.1

suppressMessages(library(tidyr))

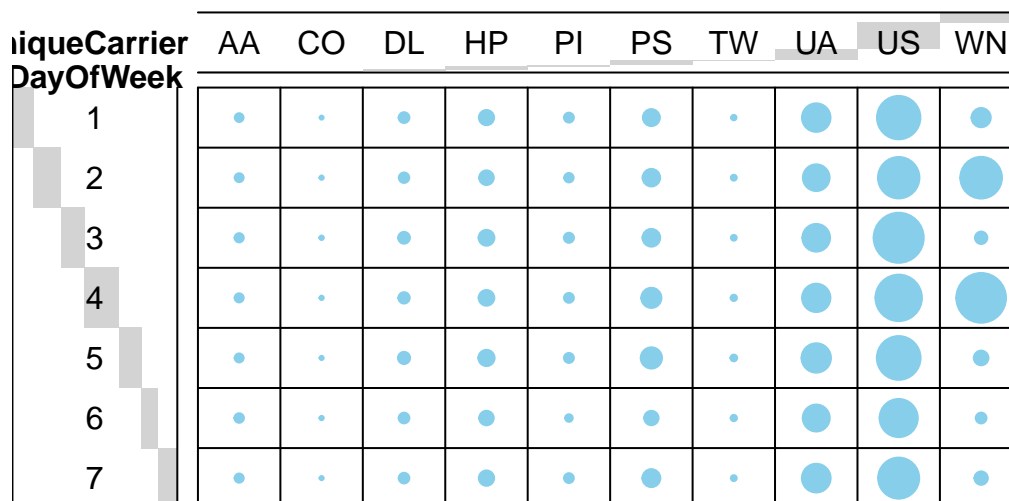
## Warning: package 'tidyr' was built under R version 3.3.2

suppressMessages(library(dplyr))
## Checks on departure delays
data.filter4 <- data %>% filter(IsDepDelayed == 'YES')
data.select <- data.filter4 %>% select(DayOfWeek,UniqueCarrier)
tab<- (table(data.select))
tab

##      UniqueCarrier
## DayOfWeek AA CO DL HP PI PS TW UA US WN
##      1 58 12 88 175 74 219 22 617 1427 280
##      2 60 14 83 162 67 234 26 540 1303 1315
##      3 64 15 105 190 77 238 25 581 1897 117
##      4 62 14 98 187 74 307 30 607 1634 1870
##      5 61 13 112 200 75 341 34 653 1448 163
##      6 69 14 87 160 43 157 30 552 1103 88
##      7 63 13 86 178 73 244 26 625 1276 136
```

```
# Plot a graphical matrix
balloonplot(t(tab), main ="Funny Balloon Plot", ylab ="DayOfWeek", xlab="UniqueCarrier",
            label = FALSE, show.margins = FALSE)
```

Funny Balloon Plot



Boxplots

Checks on Departures Delayed from Boston between Year 2000 and 2006

```
# Library
library(ggplot2)
library(tidyr)
library(dplyr)
library(shiny)
library(ggvis)
library(corrplot)
library(data.table)
library(lubridate)

require(Matrix)
if (!require('vcd')) install.packages('vcd')
library(Ckmeans.1d.dp)

## Checks on Departures Delayed from Boston between Year 2000 and 2006

data.filter1 <- data %>% filter(IsDepDelayed == 'YES', Year >= 2000 & Year<2007,
                                Origin == 'BOS')

par(las=2,                                # use perpendicular axis labels
```




```

mar=c(10.1,4.1,4.1,2.1),      # create enough space for long x labels
mgp=c(8,1,0)                   # move x axis legend down to avoid overlap
)

# start with simple boxplot with options
boxplot(DepDelay~UniqueCarrier,data=data.filter1)

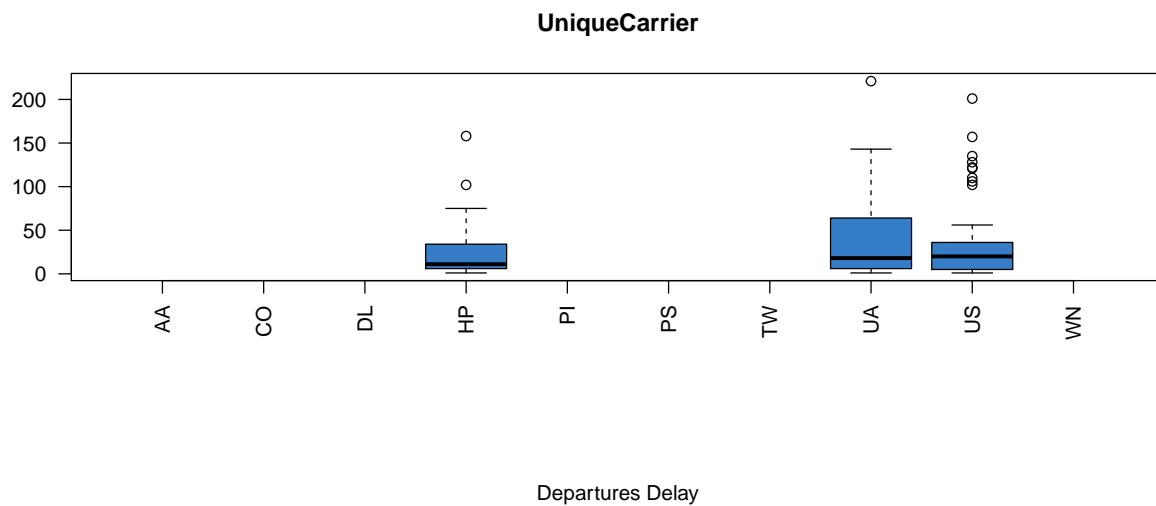
```



```

#par(mfrow=c(1,3))
boxplot(DepDelay~UniqueCarrier,data=data.filter1, main="UniqueCarrier", xlab="Departures Delay",
        ylab="Lifetime",col="#357EC7")


```



```


boxplot(DepDelay~DayOfWeek,data=data.filter1, main="Day Of Week", xlab="Departures Delay",
        ylab="",col="#357EC7") # Day Of Week = 1 is Monday

```



Departures Delay

```
boxplot(DepDelay~UniqueCarrier+DayOfWeek,data=data.filter1,
        main="UniqueCarrier & DayOfWeek", xlab="Departures Delay", ylab="Lifetime",col="#357EC7")
```




Departures Delay

Histogram

```
## Histogram of Distance on full data set
with(data,hist(Distance))
```

Histogram of Distance



```
## Checks on early arrivals from Boston
data.filter2 <- data %>% filter(IsDepDelayed == 'NO', ArrDelay == 0, Origin == 'BOS')
summary(data.filter2)
```

```
##      Year      Month   DayOfMonth   DayOfWeek   DepTime
##  Min.   :1991   Min.    :1      Min.    : 4.00   Min.     :1      Min.    : 559.0
## 1st Qu.:1998   1st Qu.:1      1st Qu.: 9.75   1st Qu.:2      1st Qu.: 944.5
## Median :2002   Median :1      Median :12.00   Median :3      Median :1328.0
## Mean   :2001   Mean    :1      Mean    :13.90   Mean    :3      Mean    :1294.0
## 3rd Qu.:2004   3rd Qu.:1      3rd Qu.:15.50   3rd Qu.:4      3rd Qu.:1656.5
## Max.    :2006   Max.     :1      Max.     :31.00   Max.     :6      Max.     :1959.0
```

```
##
##      CRSDepTime      ArrTime      CRSArrTime   UniqueCarrier
##  Min.    : 605      Min.    : 750      Min.    : 750      US           :11
## 1st Qu.: 945      1st Qu.:1294      1st Qu.:1294      HP           : 5
## Median :1338      Median :1549      Median :1549      UA           : 4
## Mean    :1310      Mean    :1527      Mean    :1527      AA           : 0
## 3rd Qu.:1661      3rd Qu.:1876      3rd Qu.:1876      CO           : 0
## Max.    :2005      Max.    :2138      Max.    :2138      DL           : 0
##                                     (Other): 0
```

```
##      FlightNum      TailNum   ActualElapsedTime CRSElapsedTime
##  Min.    : 101.0      N815AA : 2      Min.    : 70.0      Min.    : 67.0
## 1st Qu.: 425.5      N300AA : 1      1st Qu.: 97.5      1st Qu.: 92.0
## Median : 759.0      N327UA : 1      Median :116.0      Median :115.0
## Mean    :1002.9      N355US : 1      Mean    :179.1      Mean    :175.0
## 3rd Qu.:1317.8      N371AA : 1      3rd Qu.:295.2      3rd Qu.:286.2
```


```

## Max. :2792.0 (Other):11 Max. :351.0 Max. :348.0
## NA's : 3
## AirTime ArrDelay DepDelay Origin Dest
## Min. : 44.0 Min. :0 Min. : -12.00 BOS :20 PHX :5
## 1st Qu.: 73.0 1st Qu.:0 1st Qu.: -6.00 ABE : 0 PIT :4
## Median :104.0 Median :0 Median : -4.00 ABQ : 0 PHL :3
## Mean :165.8 Mean :0 Mean : -4.05 ACY : 0 ORD :2
## 3rd Qu.:311.0 3rd Qu.:0 3rd Qu.: -1.75 ALB : 0 ROC :2
## Max. :326.0 Max. :0 Max. : 0.00 AMA : 0 CLE :1
## NA's :3 (Other): 0 (Other):3
## Distance TaxiIn TaxiOut Cancelled
## Min. : 185 Min. : 3.000 Min. : 5.00 Min. :0
## 1st Qu.: 343 1st Qu.: 4.000 1st Qu.:16.00 1st Qu.:0
## Median : 496 Median : 7.000 Median :16.00 Median :0
## Mean : 983 Mean : 7.941 Mean :18.35 Mean :0
## 3rd Qu.:1890 3rd Qu.:11.000 3rd Qu.:21.00 3rd Qu.:0
## Max. :2300 Max. :21.000 Max. :40.00 Max. :0
## NA's :3 NA's :3
## CancellationCode Diverted CarrierDelay WeatherDelay NASDelay
## : 5 Min. :0 Min. :0 Min. :0 Min. :0
## A : 0 1st Qu.:0 1st Qu.:0 1st Qu.:0 1st Qu.:0
## B : 0 Median :0 Median :0 Median :0 Median :0
## C : 0 Mean :0 Mean :0 Mean :0 Mean :0
## NA's:15 3rd Qu.:0 3rd Qu.:0 3rd Qu.:0 3rd Qu.:0
## Max. :0 Max. :0 Max. :0 Max. :0
## NA's :15 NA's :15 NA's :15
## SecurityDelay LateAircraftDelay IsArrDelayed IsDepDelayed groupDepDelay
## Min. :0 Min. :0 NO :20 NO :20 (-16,-2]:15
## 1st Qu.:0 1st Qu.:0 YES: 0 YES: 0 (-2,1] : 5
## Median :0 Median :0 (1,10] : 0
## Mean :0 Mean :0 (10,473]: 0
## 3rd Qu.:0 3rd Qu.:0
## Max. :0 Max. :0
## NA's :15 NA's :15

## Histogram of Distance on filter dataset
with(data.filter2,hist(Distance))

```


Histogram of Distance



Visualization of a correlation matrix


```
## Filter and List only numeric columns from a data frame
nums <- sapply(data, is.numeric)
data.filter3 <- data[,nums]

# Replace NAs in data with 0
data.filter3[is.na(data.filter3)] <- 0
cor.matrix <- cor(data.filter3, use = "na.or.complete")
corrplot(cor.matrix)
```



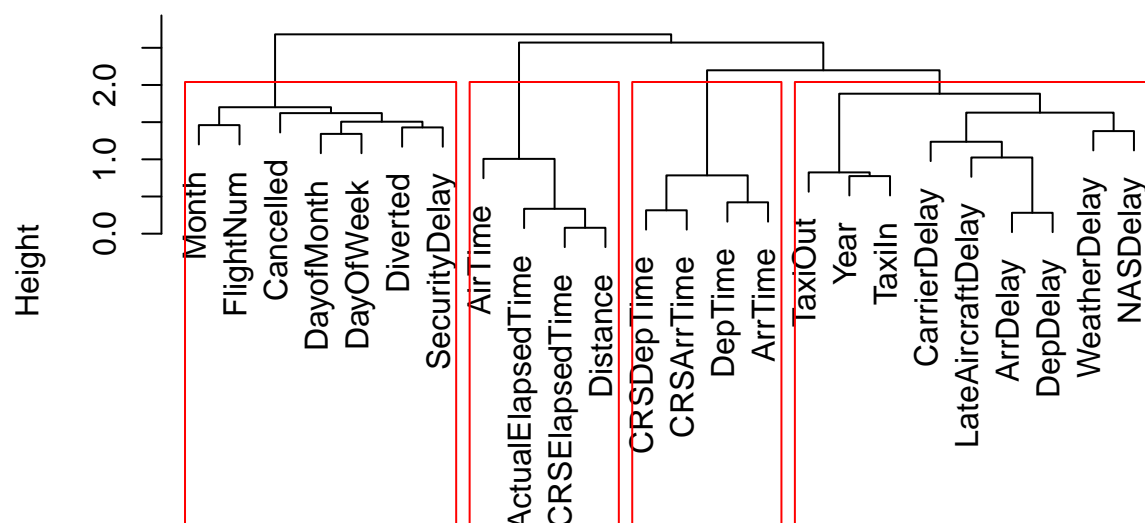
?corrplot # for more details

```
corrplot(cor.matrix, order="hclust", hclust.method="complete", addrect = 4)
```



```
# names(cor.matrix) <- colnames(data.filter3)
cor.clust <- hclust(dist(cor.matrix), method = "complete")
plot(cor.clust)
rect.hclust(cor.clust, k = 4)
```


Cluster Dendrogram



```
dist(cor.matrix)
hclust (*, "complete")
```

Time Series Analysis

```
library(quantmod)
library(fBasics)
library(tseries)
library(Hmisc)
```

```
getSymbols("~FCHI",src="yahoo")
```


```
## [1] "FCHI"
```

```
chartSeries(FCHI) # CAC 40
```



```
write.zoo(FCHI,"FCHI.csv",index.name="Date",sep=",") # save an xts object to csv
```

```
FCHI.import<- read.csv('FCHI.csv')
FCHI <- as.timeSeries(FCHI.import)
plot(FCHI, main="CAC 40")
```



Checking Basic Statistics


```
basicStats(FCHI) # Basic Statistics
```

##	FCHI.Open	FCHI.High	FCHI.Low	FCHI.Close
## nobs	2.591000e+03	2.591000e+03	2.591000e+03	2.591000e+03
## NAs	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
## Minimum	2.519430e+03	2.552990e+03	2.465460e+03	2.519290e+03
## Maximum	6.157330e+03	6.168150e+03	6.112290e+03	6.168150e+03
## 1. Quartile	3.632895e+03	3.660780e+03	3.591075e+03	3.625420e+03

```
## 3. Quartile 4.572955e+03 4.599495e+03 4.533880e+03 4.565015e+03
## Mean 4.171737e+03 4.202615e+03 4.135949e+03 4.170186e+03
## Median 4.110850e+03 4.136610e+03 4.080150e+03 4.108270e+03
## Sum 1.080897e+07 1.088897e+07 1.071624e+07 1.080495e+07
## SE Mean 1.512094e+01 1.510276e+01 1.514308e+01 1.512350e+01
## LCL Mean 4.142087e+03 4.173000e+03 4.106255e+03 4.140531e+03
## UCL Mean 4.201387e+03 4.232229e+03 4.165643e+03 4.199842e+03
## Variance 5.924133e+05 5.909896e+05 5.941499e+05 5.926142e+05
## Stdev 7.696839e+02 7.687585e+02 7.708112e+02 7.698144e+02
## Skewness 4.860190e-01 4.953230e-01 4.721540e-01 4.859360e-01
## Kurtosis -2.516240e-01 -2.599530e-01 -2.459210e-01 -2.510850e-01
## FCHI.Volume FCHI.Adjusted
## nobs 2.591000e+03 2.591000e+03
## NAs 0.000000e+00 0.000000e+00
## Minimum 0.000000e+00 2.519290e+03
## Maximum 5.312476e+08 6.168150e+03
## 1. Quartile 1.011688e+08 3.625420e+03
## 3. Quartile 1.571855e+08 4.565015e+03
## Mean 1.321279e+08 4.170186e+03
## Median 1.276144e+08 4.108270e+03
## Sum 3.423435e+11 1.080495e+07
## SE Mean 1.092376e+06 1.512350e+01
## LCL Mean 1.299859e+08 4.140531e+03
## UCL Mean 1.342699e+08 4.199842e+03
## Variance 3.091802e+15 5.926142e+05
## Stdev 5.560398e+07 7.698144e+02
## Skewness 9.072960e-01 4.859360e-01
## Kurtosis 4.371497e+00 -2.510850e-01
```


Visualise the Closing Price of CAC 40

```
cac40.close <- as.timeSeries(FCHI[,4]) #cac40.close <- as.ts(FCHI[,4],start=c(2007,01,02))
plot(cac40.close, type = "l", col = "steelblue", main = "CAC 40")
abline(h = 0, col = "grey")
```

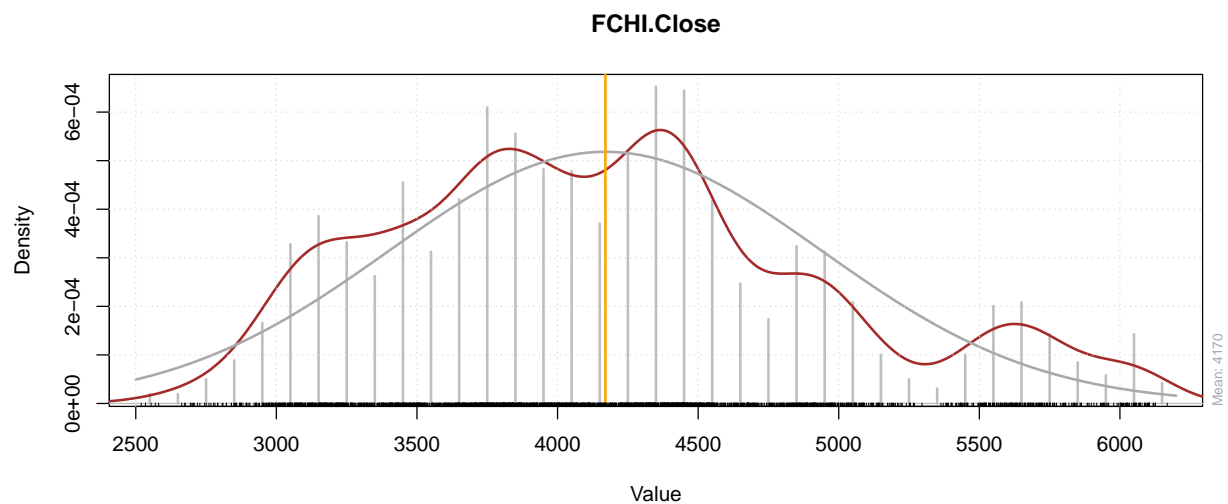


Histogram & Density Plots:

```
histPlot(cac40.close)
```



```
densityPlot(cac40.close)
```



Descriptive Statistics on the Closing Price of CAC 40

```
summary(cac40.close)
```

```
##    FCHI.Close  
##  Min.   :2519  
## 1st Qu.:3625  
## Median :4108  
## Mean   :4170  
## 3rd Qu.:4565
```

```
## Max. :6168
```

```
describe(cac40.close)
```

```
## cac40.close [FCHI.Close] Format:%Y-%m-%d
##      n missing distinct      Info      Mean      Gmd      .05      .10
##  2591      0      2577        1      4170     867.2     3052     3186
##    .25     .50     .75     .90     .95
##  3625    4108    4565    5381    5682
##
## lowest : 2519.29 2534.45 2554.55 2569.63 2581.46
## highest: 6117.96 6120.20 6125.60 6125.81 6168.15
```

Hypothesis Testing

Tests the null of normality:

```
jarque.bera.test(cac40.close) # Tests the null of normality
```

```
##
## Jarque Bera Test
##
## data:  cac40.close
## X-squared = 108.78, df = 2, p-value < 2.2e-16
```

Test for Trend Stationarity

```
kpss.test(cac40.close, null = "Trend") # KPSS Test for Trend Stationarity
```

```
##
## KPSS Test for Trend Stationarity
##
## data:  cac40.close
## KPSS Trend = 3.433, Truncation lag parameter = 11, p-value = 0.01
```

Augmented Dickey-Fuller Test for Stationarity:

```
tseries::adf.test(cac40.close, k = 10) # Augmented Dickey-Fuller Test
```

```
##
## Augmented Dickey-Fuller Test
##
## data:  cac40.close
## Dickey-Fuller = -1.9931, Lag order = 10, p-value = 0.5813
## alternative hypothesis: stationary
```

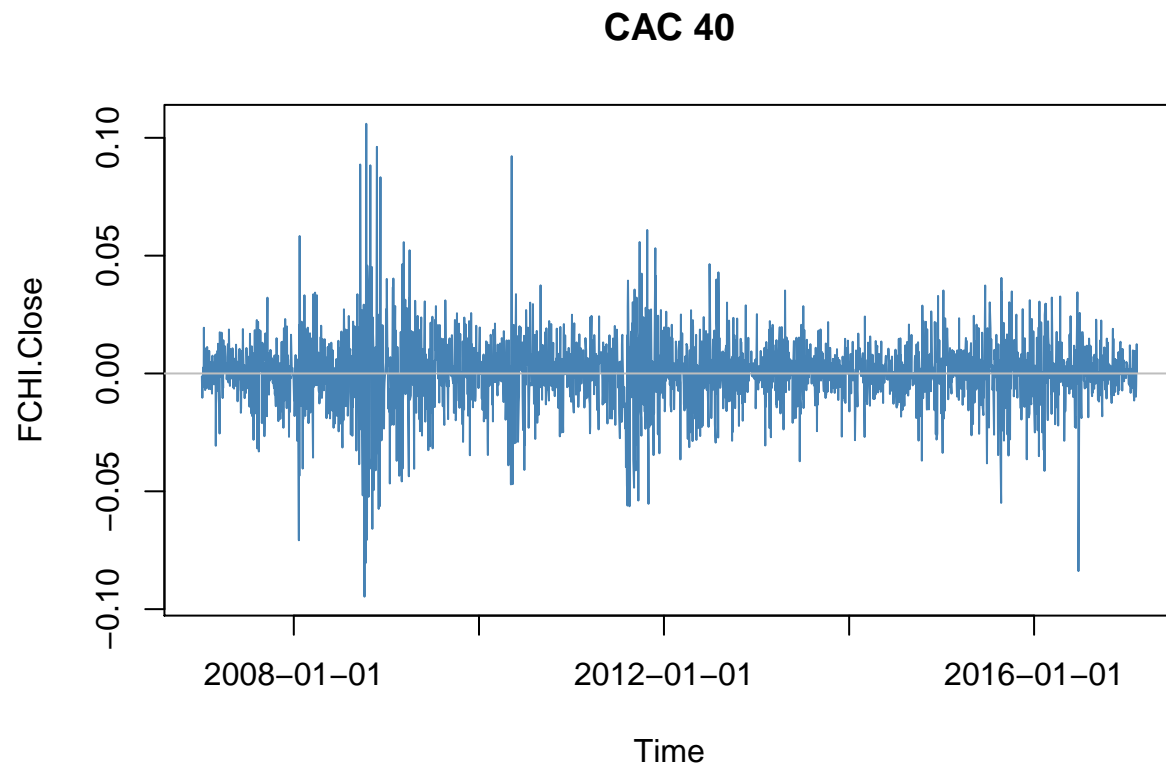
```
adf.test(diff(log(cac40.close))[-1,], alternative="stationary", k=0) # Augmented Dickey-Fuller Test
```

```
##
## Augmented Dickey-Fuller Test
##
## data:  diff(log(cac40.close))[-1, ]
```

```
## Dickey-Fuller = -52.803, Lag order = 0, p-value = 0.01  
## alternative hypothesis: stationary
```

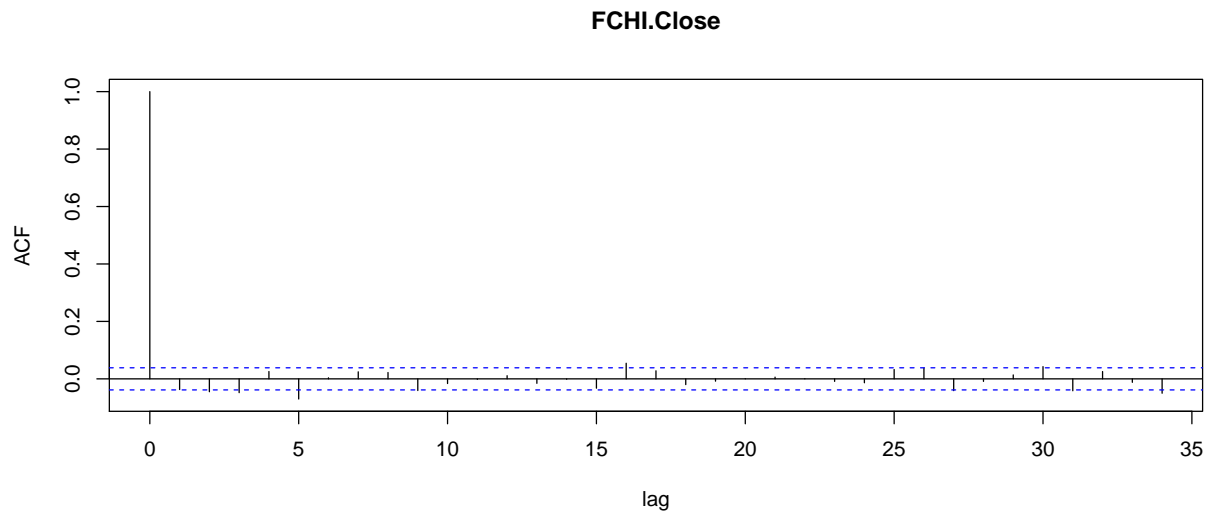
Time Series of difference of log CAC 40

```
cac40.close.df <- diff(log(cac40.close))[-1,]  
plot(cac40.close.df, type = "l", col = "steelblue", main = "CAC 40")  
abline(h = 0, col = "grey")
```




Autocorrelation Function Plots

```
acfPlot(cac40.close.df)
```



```
pacfPlot(cac40.close.df)
```



Modeling

```
fit <- arima(cac40.close.df, order = c(5, 0, 0))
fit
```

```
##
## Call:
## arima(x = cac40.close.df, order = c(5, 0, 0))
##
## Coefficients:
##          ar1      ar2      ar3      ar4      ar5  intercept
##      -0.0388 -0.0511 -0.0544  0.0167 -0.0733      -1e-04
## s.e.   0.0196   0.0196   0.0196  0.0196  0.0196       2e-04
##
```




```
## sigma^2 estimated as 0.0002277: log likelihood = 7186.52, aic = -14359.04
```

```
summary(fit)
```


```
##           Length Class  Mode
## coef           6  -none- numeric
## sigma2          1  -none- numeric
## var.coef        36 -none- numeric
## mask            6  -none- logical
## loglik           1  -none- numeric
## aic              1  -none- numeric
## arma             7  -none- numeric
## residuals 2590    ts      numeric
## call            3  -none- call
## series           1  -none- character
## code             1  -none- numeric
## n.cond           1  -none- numeric
## nobs             1  -none- numeric
## model            10 -none- list
```

```
tsdiag(fit)
```


Standardized Residuals



ACF of Residuals



p values for Ljung-Box statistic



Introduction to Data Science Algorithms in R

Ensemble Feature Selection on Titanic Dataset

```
# install.packages("fscaret", dependencies = c("Depends", "Suggests"))

library(fscaret)
library(caret)

# list of models on fscaret packages:
data(funcRegPred)
model.list <- data.frame(models=funcRegPred)

# list of models on caret packages:
names(getModelInfo())

# Import dataset from the UCI Machine Learning Repository
# (http://archive.ics.uci.edu/ml/datasets.html)
# titanic.data <- read.csv('http://math.ucdenver.edu/RTutorial/titanic.txt', sep='\t')
titanic.data <- read.csv('titanic.txt', sep='\t') # import data from working directory

# creating new title feature
titanic.data$Title <- ifelse(grepl('Mr ', titanic.data$Name), 'Mr',
                             ifelse(grepl('Mrs ', titanic.data$Name), 'Mrs',
                                     ifelse(grepl('Miss', titanic.data$Name), 'Miss', 'Nothing')))
titanic.data$Title <- as.factor(titanic.data$Title)

# Replace NAs in Age with the median age
titanic.data$Age[is.na(titanic.data$Age)] <- median(titanic.data$Age, na.rm=T)

# reorder data set so response variable is last column
titanic.data <- titanic.data[c('PClass', 'Age', 'Sex', 'Title', 'Survived')]

# binarize all factors
titanicDummy <- dummyVars("~.", data=titanic.data, fullRank=F)
titanic.data <- as.data.frame(predict(titanicDummy, titanic.data))
```

We have a look at the structure of data

```
str(titanic.data)

## 'data.frame':    1313 obs. of  11 variables:
## $ PClass.1st    : num  1 1 1 1 1 1 1 1 1 1 ...
## $ PClass.2nd    : num  0 0 0 0 0 0 0 0 0 0 ...
## $ PClass.3rd    : num  0 0 0 0 0 0 0 0 0 0 ...
## $ Age           : num  29 2 30 25 0.92 47 63 39 58 71 ...
## $ Sex.female    : num  1 1 0 1 0 0 1 0 1 0 ...
## $ Sex.male      : num  0 0 1 0 1 1 0 1 0 1 ...
## $ Title.Miss    : num  1 1 0 0 0 0 1 0 0 0 ...
## $ Title.Mr      : num  0 0 1 0 0 1 0 1 0 1 ...
## $ Title.Mrs     : num  0 0 0 1 0 0 0 0 1 0 ...
## $ Title.Nothing: num  0 0 0 0 1 0 0 0 0 0 ...
## $ Survived      : num  1 0 0 0 1 1 1 0 1 0 ...

# split data set into train and test portion
set.seed(1234)
```

```
splitIndex <- createDataPartition(titanic.data$Survived, p = .75, list = FALSE, times = 1)
# splitIndex <- sample(nrow(titanic.data), floor(0.75*nrow(titanic.data)))
trainDF <- titanic.data[ splitIndex,]
testDF <- titanic.data[-splitIndex,]
```

We now run an ensemble feature selection specifying list of models available on 'fscaret' package:

```
# limit models to use in ensemble and run fscaret
featureSelection.models <- c("glm", "gbm", "treebag", "ridge", "lasso")
featureSelection<-fscaret(trainDF, testDF, myTimeLimit = 40, preprocessData=TRUE,
                          Used.funcRegPred = featureSelection.models, with.labels=TRUE,
                          suppress.output=FALSE, no.cores=2)
```

```
# analyze results
```

```
print(featureSelection$VarImp)
```

```
## $rawMSE
##          gbm          glm          lasso          ridge          treebag
## 1 0.1286814 0.1386274 0.1386274 0.1386278 0.1298043
##
## $rawRMSE
##          gbm          glm          lasso          ridge          treebag
## 1 0.3587219 0.372327 0.372327 0.3723276 0.3602836
##
## $matrixVarImp.RMSE
##          gbm glm          lasso          ridge          treebag          SUM          SUM%
## 5 28.1448932 0 25.0403181 25.0400842 13.307300 91.532596 100.000000
## 7 27.2664317 0 24.2347615 24.2353151 13.670037 89.406545 97.677275
## 3 26.0342484 0 15.2132114 15.2127921 24.635434 81.095686 88.597604
## 1 5.4238748 0 11.2699282 11.2694387 13.467735 41.430977 45.263631
## 8 0.4340408 0 12.1964247 12.1966233 6.545685 31.372774 34.274974
## 4 11.2236464 0 0.1074810 0.1075533 18.074952 29.513633 32.243850
## 6 0.8201142 0 6.3783194 6.3784115 2.427869 16.004714 17.485262
## 2 0.0000000 0 1.4203760 1.4204060 4.552760 7.393542 8.077497
## 9 0.6527505 0 0.4850858 0.4851359 2.884740 4.507713 4.924708
##          ImpGrad Input_no
## 5 0.000000 5
## 7 2.322725 7
## 3 9.295582 3
## 1 48.910998 1
## 8 24.277012 8
## 4 5.925968 4
## 6 45.771792 6
## 2 53.803973 2
## 9 39.031759 9
##
## $matrixVarImp.MSE
##          gbm glm          lasso          ridge          treebag          SUM          SUM%
## 5 28.1448932 0 24.1253213 24.1250595 13.249615 89.644889 100.000000
## 7 27.2664317 0 23.3492005 23.3496986 13.610779 87.576110 97.692251
## 3 26.0342484 0 14.6573064 14.6568801 24.528643 79.877078 89.103884
## 1 5.4238748 0 10.8581144 10.8576264 13.409354 40.548970 45.232885
## 8 0.4340408 0 11.7507559 11.7509294 6.517310 30.453036 33.970745
## 4 11.2236464 0 0.1035535 0.1036230 17.996600 29.427423 32.826660
## 6 0.8201142 0 6.1452496 6.1453290 2.417345 15.528038 17.321721
```

```
## 2 0.0000000 0 1.3684741 1.3685009 4.533025 7.270000 8.109776
## 9 0.6527505 0 0.4673603 0.4674078 2.872236 4.459754 4.974912
##      ImpGrad Input_no
## 5 0.0000000      5
## 7 2.307749      7
## 3 8.791247      3
## 1 49.235787      1
## 8 24.898126      8
## 4 3.367852      4
## 6 47.232764      6
## 2 53.181465      2
## 9 38.655374      9
##
## $model
## list()
```

```
print(featureSelection$PPlabels)
```

```
##      Orig Input No      Labels
## 1          1      PClass.1st
## 2          2      PClass.2nd
## 3          3      PClass.3rd
## 4          4          Age
## 5          6      Sex.male
## 6          7      Title.Miss
## 7          8      Title.Mr
## 8          9      Title.Mrs
## 9         10 Title.Nothing
```

References (Some useful R links)

I do recommend that you simply use search engines to find out the links to resources that suits you. Here are a few of such links:

Introduction:

- Introduction to R <https://cran.r-project.org/doc/manuals/R-intro.html>
- Interactive intro to R programming language <https://www.datacamp.com/courses/introduction-to-r>
- for data manipulation <https://cran.rstudio.com/web/packages/dplyr/vignettes/introduction.html>

Graphics:

- Tutorial on plots <http://cran.r-project.org/doc/contrib/Rossiter-RIntro-ITC.pdf>

Statistics:

- Quick-R <http://www.statmethods.net/>
- Beginner's tutorial for Time Series <http://www.stat.pitt.edu/stoffer/tsa2/index.html>

Last updated on February 2017, SNCF Division MMI